

TIRUMALARAO KILARI

AI Agent Developer · Backend Engineer · LLM Systems Builder

346-474-6746 · kilaritirumalarao@gmail.com · [linkedin.com/in/tirumalaraokilari-803829273](https://www.linkedin.com/in/tirumalaraokilari-803829273) · github.com/bizygoodz-web · Pflugerville, TX

PROFESSIONAL SUMMARY

AI Engineer with hands-on experience designing and evaluating production-grade agentic systems — including multi-stage pipelines, multi-turn LLM interactions, and modular backend architectures. Built a fully autonomous RAG agent featuring session state management, prompt injection safeguards, and a RAGAS-style evaluation harness. Shipped ResumeAI — a live web application that automatically reads job postings, tailors resumes using LLM prompt engineering, and generates cover letters. Experienced in Python, JavaScript, and SQL across live, non-mocked environments.

TECHNICAL SKILLS

Agent & LLM: LangGraph, LangChain, Anthropic Claude, OpenAI GPT-4, Google Gemini, Prompt Engineering

Backend & APIs: FastAPI, Uvicorn, REST APIs, Pydantic, Python (primary), JavaScript, SQL, Bash

RAG & Retrieval: FAISS, Sentence-Transformers, Cosine Similarity, Chunking Strategies, Embedding Pipelines

Web & Deployment: Streamlit, Render, Railway, GitHub Actions (CI/CD), Git, GitHub, httpx, BeautifulSoup

Security & Eval: Prompt injection detection, RAGAS-style eval harness, faithfulness scoring, hallucination prevention

WORK EXPERIENCE

AI Agent & Prompt Engineer (Contract) | **Blue Origin IT Staffing — Client: Austinite Market** · 2024 – Present · Austin, TX

- Designed and deployed a multi-stage AI content pipeline: client brief ingestion → prompt engineering → image generation → iterative refinement → production delivery
- Engineered detailed system prompts in Google Gemini to autonomously produce brand-consistent marketing assets; iterated across multiple rounds adjusting style, composition, and output quality
- Delivered production-ready assets physically deployed at the client storefront — validating real-world AI output quality in a live, non-mocked environment
- Demonstrated ability to translate complex business requirements into precise AI agent instructions — core to training LLMs for autonomous, multi-step task execution

PROJECTS

ResumeAI — Live Web Application · Python · Streamlit · Groq · BeautifulSoup · Render · resumeai-v2.onrender.com

- Built and deployed a full-stack AI web application that reads any job posting URL automatically, tailors resume bullet points using LLM prompt engineering, generates a personalized cover letter, and pre-fills answers to common application questions
- Engineered a multi-stage pipeline: URL scraping → job text extraction → skill gap analysis → resume rewriting → cover letter generation → application package download — zero manual input after URL paste
- Deployed live on Render with Streamlit frontend; handles PDF/DOCX resume parsing, real-time job scraping with BeautifulSoup, and Groq LLaMA API calls

AI Agent with Web Search — Live Web Application · Python · Streamlit · Groq · DuckDuckGo · my-ai-agent-app.onrender.com

- Built and deployed a conversational AI agent with real-time web search capability using Groq LLaMA and DuckDuckGo API — live at my-ai-agent-app.onrender.com

- Implemented a tool-use agent loop: LLM decides whether to search the web, executes the search, feeds results back to the model, and generates a grounded final answer
- Full multi-turn conversation memory with streaming responses and live search status indicators

RAG PDF Chat Agent · Python · FAISS · Claude · FastAPI · pytest · github.com/bizygoodz-web/rag-pdf-chat

- Built a production-grade autonomous RAG agent with full modular separation: distinct services for PDF ingestion, vector indexing, semantic retrieval, and Claude-powered response generation
- Implemented persistent session state via FAISS index saving/reloading and prompt injection defenses — agent refuses out-of-context requests, preventing authority escalation
- Shipped FastAPI REST server with /ingest, /query, /health endpoints; 10 unit tests (all passing) with GitHub Actions CI/CD; RAGAS-style eval harness measuring citation rate and faithfulness

EDUCATION

Computer Science | [Lindsey Wilson College](#) · 2022 – 2023 · Columbia, KY

- Independently built RAG systems, LLM pipelines, and agentic workflows; self-directed study in LangGraph, RAGAS evaluation, and Anthropic prompt engineering best practices

CERTIFICATIONS & CONTINUOUS LEARNING

- DeepLearning.AI — LangChain & Vector Databases in Production (2025, in progress)
- Daily use of GitHub Copilot for AI-assisted development, code review, and agentic workflow design
- Deep hands-on experience with Claude API (claude-sonnet-4) — system prompts, token management, grounding strategies, and identifying subtle LLM failure modes
- Studying: LangGraph agentic workflows, persistent agent state management, and multi-agent coordination patterns

Open to remote and US-based roles · Available immediately · github.com/bizygoodz-web